



VictoriaMetrics / vmaalert

# vmaalert



`vmaalert` executes a list of the given [alerting](#) or [recording](#) rules against configured `-datasource.url`. For sending alerting notifications `vmaalert` relies on [Alertmanager](#) configured via `-notifier.url` flag. Recording rules results are persisted via [remote write](#) protocol and require `-remoteWrite.url` to be configured. `vmaalert` is heavily inspired by [Prometheus](#) implementation and aims to be compatible with its syntax.

A [single-node](#) or [cluster version](#) of VictoriaMetrics are capable of proxying requests to `vmaalert` via `-vmaalert.proxyURL` command-line flag. Use this feature for the following cases:

- for proxying requests from [Grafana Alerting UI](#) ;
- for accessing `vmaalert` 's UI through VictoriaMetrics Web interface.

[VictoriaMetrics Cloud](#) provides out-of-the-box alerting functionality based on `vmaalert` . This service simplifies the setup and management of alerting and recording rules as well as the integration with Alertmanager. For more details, please refer to the [VictoriaMetrics Cloud documentation](#) .

## Features

- Integration with [VictoriaMetrics](#) and [MetricsQL](#) ;
- Integration with [VictoriaLogs](#) and [LogsQL](#) . See [this doc](#) ;
- Prometheus [alerting rules definition format](#) support;
- Integration with [Alertmanager](#) starting from [Alertmanager v0.16.0-alpha](#) ;
- Keeps the alerts [state on restarts](#) ;
- Graphite datasource can be used for alerting and recording rules. See [these docs](#) ;
- Recording and Alerting rules backfilling (aka `replay` ). See [these docs](#) ;
- Lightweight and without extra dependencies.
- Supports [reusable templates](#) for annotations;
- Load of recording and alerting rules from local filesystem, URL, GCS and S3;
- Detect alerting rules which [don't match any series](#) .

## Limitations

- `vmaalert` execute queries against remote datasource which has reliability risks because of the network. It is recommended to configure alerts thresholds and rules expressions with the understanding that network requests may fail;
- `vmaalert` executes rules within a group sequentially, but persistence of execution results to remote storage is asynchronous. Hence, user shouldn't rely on chaining of recording rules when result of previous recording rule is reused in the next one. See how to chain groups [here](#) .

## QuickStart

To start using `vmaalert` you will need the following things:

- list of rules - PromQL/MetricsQL expressions to execute;
- datasource address - reachable endpoint with [Prometheus HTTP API](#) support for running queries against;

- notifier address [optional] - reachable [Alert Manager](#) instance for processing, aggregating alerts, and sending notifications. Please note, notifier address also supports Consul and DNS Service Discovery via [config file](#).
- remote write address [optional] - [remote write](#) compatible storage to persist rules and alerts state info. To persist results to multiple destinations use vmagent configured with multiple remote writes as a proxy;
- remote read address [optional] - MetricsQL compatible datasource to restore alerts state from.

You can use the existing [docker-compose environment](#) as example. It already contains vmaalert configured with list of alerting rules and integrated with Alert Manager and VictoriaMetrics.

Alternatively, build vmaalert from sources:



```
git clone https://github.com/VictoriaMetrics/VictoriaMetrics
cd VictoriaMetrics
make vmaalert
```

Then run vmaalert :



```
./bin/vmaalert -rule=alert.rules \           # Path to the file with rules configuration. Supports wildcard
  -datasource.url=http://localhost:8428 \    # Prometheus HTTP API compatible datasource
  -notifier.url=http://localhost:9093 \      # AlertManager URL (required if alerting rules are used)
  -notifier.url=http://127.0.0.1:9093 \      # AlertManager replica URL
  -remoteWrite.url=http://localhost:8428 \    # Remote write compatible storage to persist rules and alerts state info (required if recording rules are used)
  -remoteRead.url=http://localhost:8428 \    # MetricsQL compatible datasource to restore alerts state from
  -external.label=cluster=east-1 \          # External label to be applied for each rule
  -external.label=replica=a                 # Multiple external labels may be set
```

#### Note

To validate the syntax of configured rules simply run vmaalert with `-rule` and `-dryRun` cmd-line flags.

Note there's a separate `-remoteWrite.url` command-line flag to allow writing results of alerting/recording rules into a different storage than the initial data that's queried. This allows using vmaalert to aggregate data from a short-term, high-frequency, high-cardinality storage into a long-term storage with decreased cardinality and a bigger interval between samples. See also [stream aggregation](#).

See the full list of configuration flags in [configuration](#) section.

If you run multiple vmaalert services for the same datastore or AlertManager - do not forget to specify different `-external.label` command-line flags in order to define which vmaalert generated rules or alerts. If rule result metrics have label that conflict with `-external.label`, vmaalert will automatically rename it with prefix `exported_`.

Configuration for [recording](#) and [alerting](#) rules is very similar to Prometheus rules and configured using YAML. Configuration examples may be found in [testdata](#) folder. Every `rule` belongs to a `group` and every configuration file may contain arbitrary number of groups:



```
groups:
[ - <rule_group> ]
```

#### Note

Explore how to integrate vmaalert with [VictoriaMetrics Anomaly Detection](#) in the following [guide](#).

#### Note

For users of [VictoriaMetrics Cloud](#), many of the configuration steps (including highly available setup of vmaalert for cluster deployments) are handled automatically. Please, refer to the [VictoriaMetrics Cloud documentation](#) for more details.

## Groups

Each group has the following attributes:



```
# The name of the group. Must be unique within a file.
name: <string>

# How often rules in the group are evaluated.
[ interval: <duration> | default = -evaluationInterval flag ]

# Optional
# Group will be evaluated at the exact offset in the range of [0...interval].
# E.g. for Group with `interval: 1h` and `eval_offset: 5m` the evaluation will
# start at 5th minute of the hour. See https://github.com/VictoriaMetrics/VictoriaMetrics/issues/3409
# `interval` must be specified if `eval_offset` is used, and `eval_offset` cannot exceed `interval`.
# `eval_offset` cannot be used with `eval_delay`, as group will be executed at the exact offset and `eval_delay` is ignored.
[ eval_offset: <duration> ]

# Optional
# Adjust the `time` parameter of group evaluation requests to compensate intentional query delay from the datasource.
# By default, the value is inherited from the `-rule.evalDelay` cmd-line flag - see its description for details.
# If group has `latency_offset` set in `params`, then it is recommended to set `eval_delay` equal to `latency_offset`.
# See https://github.com/VictoriaMetrics/VictoriaMetrics/issues/5155 and https://docs.victoriametrics.com/victoriametrics/keyconcepts/#query-latency
[ eval_delay: <duration> ]

# Limit limits the number of alerts or recording results the rule within this group can produce.
# On exceeding the limit, rule will be marked with an error and all its results will be discarded.
# 0 is no limit.
[ limit: <integer> | default 0 ]

# How many rules execute at once within a group. Increasing concurrency may speed
# up group's evaluation duration (exposed via `vmalert_iteration_duration_seconds` metric).
[ concurrency: <integer> | default = 1 ]

# Optional type for expressions inside rules to override the `-rule.defaultRuleType` (default is "prometheus") cmd-line flag.
# Supported values: "graphite", "prometheus" and "vlogs" (check https://docs.victoriametrics.com/victorialogs/vmalert/ for details).
[ type: <string> ]

# Optional
# The evaluation timestamp will be aligned with group's interval,
# instead of using the actual timestamp that evaluation happens at.
#
# It is enabled by default to get more predictable results
# and to visually align with graphs plotted via Grafana or vmui.
# When comparing with raw queries, remember to use `step` equal to evaluation interval.
#
# See https://github.com/VictoriaMetrics/VictoriaMetrics/issues/5049
# Available starting from v1.95
[ eval_alignment: <bool> | default true ]

# Optional list of HTTP URL parameters
# applied for all rules requests within a group
# For example:
# params:
#   nocache: ["1"]           # disable caching for vmselect
#   denyPartialResponse: ["true"] # fail if one or more vmstorage nodes returned an error
#   extra_label: ["env=dev"]    # apply additional label filter "env=dev" for all requests
# see more details at https://docs.victoriametrics.com/victoriametrics/single-server-victoriametrics/#prometheus-querying-api-enhancements
params:
  [ <string>: [<string>, ...] ]

# Optional list of HTTP headers in form `header-name: value`
# applied for all rules requests within a group
# For example:
# headers:
#   - "CustomHeader: foo"
#   - "CustomHeader2: bar"
# Headers set via this param have priority over headers set via `-datasource.headers` flag.
headers:
  [ <string>, ... ]

# Optional list of HTTP headers in form `header-name: value`
# applied for all alert notifications sent to notifiers
# generated by rules of this group.
```

```

# It has higher priority over headers defined in notifier config.
# For example:
#   notifier_headers:
#     - "TenantID: foo"
notifier_headers:
  [ <string>, ... ]

# Optional list of labels added to every rule within a group.
# It has priority over the external labels.
# Labels are commonly used for adding environment
# or tenant-specific tag.
labels:
  [ <labelname>: <labelvalue> ... ]

rules:
  [ - <rule> ... ]

# Enable debug mode for all rules in the group.
# This can be overridden by the `debug` field in rule.
[ debug: <bool> | default = false ]

```

## Rules

Every rule contains `expr` field for [PromQL](#) or [MetricsQL](#) expression. `vmaalert` will execute the configured expression and then act according to the Rule type.

There are two types of Rules:

- [alerting](#) - Alerting rules allow defining alert conditions via `expr` field and to send notifications to [Alertmanager](#) if execution result is not empty.
- [recording](#) - Recording rules allow defining `expr` which result will be then backfilled to configured `-remoteWrite.url`. Recording rules are used to precompute frequently needed or computationally expensive expressions and save their result as a new set of time series.

`vmaalert` forbids defining duplicates - rules with the same combination of name, expression, and labels within one group.

## Alerting rules

The syntax for alerting rule is the following:



```

# The name of the alert. Must be a valid metric name.
alert: <string>

# The expression to evaluate. The expression language depends on the type value.
# By default, PromQL/MetricsQL expression is used. If group.type="graphite", then the expression
# must contain valid Graphite expression.
expr: <string>

# Alerts are considered firing once they have been returned for this long.
# Alerts which have not yet been fired for long enough are considered pending.
# If param is omitted or set to 0 then alerts will be immediately considered
# as firing once they return.
[ for: <duration> | default = 0s ]

# Alert will continue firing for this long even when the alerting expression no longer has results.
# This allows you to delay alert resolution.
[ keep_firing_for: <duration> | default = 0s ]

# Whether to print debug information into logs.
# Information includes alerts state changes and requests sent to the datasource.
# Please note, that if rule's query params contain sensitive
# information - it will be printed to logs.
# Logs are printed with INFO level, so make sure that -loggerLevel=INFO to see the output.
[ debug: <bool> | default = false ]

# Defines the number of rule's updates entries stored in memory
# and available for view on rule's Details page.
# Overrides `rule.updateEntriesLimit` value for this specific rule.
# Available starting from https://docs.victoriametrics.com/victoriametrics/changelog/#v1860

```

```
[ update_entries_limit: <integer> | default 0 ]

# Labels to add or overwrite for each alert.
# In case of conflicts, original labels are kept with prefix `exported_`.
labels:
  [ <labelname>: <tmpl_string> ]

# Annotations to add to each alert.
annotations:
  [ <labelname>: <tmpl_string> ]
```

## Templating

It is allowed to use [Go templating](#) in annotations to format data, iterate over or execute expressions. The following variables are available in templating:

Variable	Description	Example
<code>\$value</code> or <code>.Value</code>	The current alert's value. Avoid using value in labels, it may cause unexpected issues.	Number of connections is {{ \$value }}
<code>\$activeAt</code> or <code>.ActiveAt</code>	The moment of <a href="#">time</a> when alert became active ( <code>pending</code> or <code>firing</code> ).	<a "-1h")).unixmilli}}"="" (parsedurationtime="" 1h")).unixmilli}}&amp;to="{{(\$activeAt.Add" href="http://vm-grafana.com/?viewPanel=&amp;from={{(\$activeAt.Add (parseDurationTime ">http://vm-grafana.com/?viewPanel=&amp;from={{(\$activeAt.Add (parseDurationTime "1h")).UnixMilli}}&amp;to={{(\$activeAt.Add (parseDurationTime "-1h")).UnixMilli}}</a>
<code>\$labels</code> or <code>.Labels</code>	The list of labels of the current alert. Use as <code>".Labels.&lt;label_name&gt;".</code>	Too high number of connections for {{ .Labels.instance }}
<code>\$type</code> or <code>.Type</code>	The rule type: "graphite", "prometheus" or "vlogs"	Link: 'explore?left={"datasource":'{{ if eq .Type "vlogs" }}VictoriaLogs{{ else }}VictoriaMetrics{{ end }}'}'
<code>\$alertID</code> or <code>.AlertID</code>	The current alert's ID generated by vmaalert.	Link: <code>vmaalert/alert?group_id={{.GroupID}}&amp;alert_id={{.AlertID}}</code>
<code>\$groupID</code> or <code>.GroupID</code>	The current alert's group ID generated by vmaalert.	Link: <code>vmaalert/alert?group_id={{.GroupID}}&amp;alert_id={{.AlertID}}</code>
<code>\$expr</code> or <code>.Expr</code>	Alert's expression. Can be used for generating links to Grafana or other systems.	<code>/api/v1/query?query={{ \$expr queryEscape }}</code>
<code>\$for</code> or <code>.For</code>	Alert's configured for param.	Number of connections is too high for more than {{ .For }}
<code>\$externalLabels</code> or <code>.ExternalLabels</code>	List of labels configured via <code>-external.label</code> command-line flag.	Issues with {{ \$labels.instance }} (datacenter-{{ \$externalLabels.dc }}
<code>\$externalURL</code> or <code>.ExternalURL</code>	URL configured via <code>-external.url</code> command-line flag. Used for cases when vmaalert is hidden behind proxy.	Visit {{ \$externalURL }} for more details

Additionally, vmaalert provides some extra templating functions listed [here](#) and [reusable templates](#).

## Template functions

vmaalert provides the following template functions, which can be used during [templating](#):

- `args arg0 ... argN` - converts the input args into a map with `arg0`, ..., `argN` keys.
- `externalURL` - returns the value of `-external.url` command-line flag.
- `first` - returns the first result from the input query results returned by `query` function.
- `htmlEscape` - escapes special chars in input string, so it can be safely embedded as a plaintext into HTML.
- `humanize` - converts the input number into human-readable format by adding [metric prefixes](#). For example, `100000` is converted into `100K`.
- `humanize1024` - converts the input number into human-readable format with 1024 base. For example, `1024` is converted into `1ki`.
- `humanizeDuration` - converts the input number in seconds into human-readable duration.

- `humanizePercentage` - converts the input number to percentage. For example, `0.123` is converted into `12.3%`.
- `humanizeTimestamp` - converts the input unix timestamp into human-readable time.
- `jsonEscape` - JSON-encodes the input string.
- `label name` - returns the value of the label with the given `name` from the input query result.
- `match regex` - matches the input string against the provided `regex`.
- `parseDuration` - parses the input string into duration in seconds. For example, `1h` is parsed into `3600`.
- `parseDurationTime` - parses the input string into [time.Duration](#).
- `pathEscape` - escapes the input string, so it can be safely put inside path part of URL.
- `pathPrefix` - returns the path part of the `-external.url` command-line flag.
- `query` - executes the [MetricsQL](#) query against `-datasource.url` and returns the query result. For example, `{{ query "sort_desc(process_resident_memory_bytes)" | first | value }}` executes the `sort_desc(process_resident_memory_bytes)` query at `-datasource.url` and returns the first result.
- `queryEscape` - escapes the input string, so it can be safely put inside [query arg](#) part of URL.
- `quotesEscape` - escapes the input string, so it can be safely embedded into JSON string.
- `replaceAll regex repl` - replaces all the occurrences of the `regex` in input string with the `repl`.
- `safeHtml` - marks the input string as safe to use in HTML context without the need to html-escape it.
- `sortByLabel name` - sorts the input query results by the label with the given `name`.
- `stripDomain` - leaves the first part of the domain. For example, `foo.bar.baz` is converted to `foo`. The port part is left in the output string. E.g. `foo.bar:1234` is converted into `foo:1234`.
- `stripPort` - strips `port` part from `host:port` input string.
- `strvalue` - returns the metric name from the input query result.
- `title` - converts the first letters of every input word to uppercase.
- `toLowerCase` - converts all the chars in the input string to lowercase.
- `toTime` - converts the input unix timestamp to [time.Time](#).
- `toUpper` - converts all the chars in the input string to uppercase.
- `value` - returns the numeric value from the input query result.

## Reusable templates

Like in Alertmanager you can define [reusable templates](#) to share same templates across annotations. Just define the templates in a file and set the path via `-rule.templates` flag.

For example, template `grafana.filter` can be defined as following:



```
{{ define "grafana.filter" -}}
  {{- $labels := .arg0 -}}
  {{- range $name, $label := . -}}
    {{- if (ne $name "arg0") -}}
      {{- ( or (index $labels $label) "All" ) | printf "&var-%s=%s" $label -}}
    {{- end -}}
  {{- end -}}
{{- end -}}
```

And then used in annotations:



```
groups:
- name: AlertGroupName
  rules:
  - alert: AlertName
    expr: any_metric > 100
    for: 30s
    labels:
      alertname: 'Any metric is too high'
      severity: 'warning'
    annotations:
      dashboard: '{{ $externalURL }}/d/dashboard?orgId=1{{ template "grafana.filter" (args .CommonLabels "account_id" "any_label") }}'
```

The `-rule.templates` flag supports wildcards so multiple files with templates can be loaded. The content of `-rule.templates` can be also [hot reloaded](#).

## Recording rules

The syntax for recording rules is following:



```
# The name of the time series to output to. Must be a valid metric name.
record: <string>

# The expression to evaluate. The expression language depends on the type value.
# By default, MetricsQL expression is used. If group.type="graphite", then the expression
# must contain valid Graphite expression.
expr: <string>

# Labels to add or overwrite before storing the result.
# In case of conflicts, original labels are kept with prefix `exported_`.
labels:
  [ <labelname>: <labelvalue> ]

# Whether to print debug information into logs.
# Information includes requests sent to the datasources.
# information - it will be printed to logs.
# Logs are printed with INFO level, so make sure that -loggerLevel=INFO to see the output.
[ debug: <bool> | default = false ]

# Defines the number of rule's updates entries stored in memory
# and available for view on rule's Details page.
# Overrides 'rule.updateEntriesLimit' value for this specific rule.
[ update_entries_limit: <integer> | default 0 ]
```

For recording rules to work `-remoteWrite.url` must be specified.

## Alerts state on restarts

`vmaalert` holds alerts state in the memory. Restart of the `vmaalert` process will reset the state of all active alerts in the memory. To prevent `vmaalert` from losing the state on restarts configure it to persist the state to the remote database via the following flags:

- `-remoteWrite.url` - URL to VictoriaMetrics (Single) or `vminsert` (Cluster). `vmaalert` will persist alerts state to the configured address in the form of [time series](#) `ALERTS` and `ALERTS_FOR_STATE` via remote-write protocol. These time series can be queried from VictoriaMetrics just as any other time series. The state will be persisted to the configured address on each evaluation.
- `-remoteRead.url` - URL to VictoriaMetrics (Single) or `vmselect` (Cluster). `vmaalert` will try to restore alerts state from the configured address by querying time series with name `ALERTS_FOR_STATE`. The restore happens only once when `vmaalert` process starts, and only for the configured rules. Config [hot reload](#) doesn't trigger state restore.

Both flags are required for proper state restoration. Restore process may fail if time series are missing in configured `-remoteRead.url`, weren't updated in the last 1h (controlled by `-remoteRead.lookback`) or received state doesn't match current `vmaalert` rules configuration. `vmaalert` marks successfully restored rules with `restored` label in [web UI](#).

## Link to alert source

Alerting notifications sent by vmalert always contain a `source` link. By default, the link format is the following `http://<vmalert-addr>/vmalert/alert?group_id=<group_id>&alert_id=<alert_id>`. On click, it opens vmalert [web UI](#) to show the alert status and its fields.

It is possible to override the link format. For example, to make the link to [vmui](#) specify the following cmd-line flags:



```
./bin/vmalert \
  -external.url=http://<vmui-addr> \ # the hostname and port for datasource vmui
  -external.alert.source='vmui/#/?g0.expr={{.Expr|queryEscape}}' # the path built using alert expr
```

Now, all `source` links will lead to `http://<vmui-addr>/vmui/#/?g0.expr=$expr`, where `$expr` is an alerting rule expression.

The `-external.alert.source` cmd-line flag supports [templating](#) and allows using labels and extra data related to the alert. For example, see the following link to Grafana:



```
./bin/vmalert \
  -external.url=http://<grafana-addr> \ # the hostname and port for Grafana
  -external.alert.source='explore?left={"datasource":{"if eq .Type \"vlogs\" }}VictoriaLogs{{ else }}VictoriaMetrics{{ end }}","queries":[{"
```

In this example, `-external.alert.source` will lead to Grafana's Explore page with `expr` field equal to alert expression, and time range will be selected starting from `"from":{"ActiveAt.UnixMilli }}"` when alert became active. The `datasource` name is set to `VictoriaLogs` if rule's type [②](#) (`prometheus`, `vlogs` or `graphite`) is `vlogs`. Otherwise, it is set to `VictoriaMetrics`. See [how we set alert source in docker](#).

In addition to `source` link, some extra links could be added to alert's `annotations` field. See [how we use them](#) to link alerting rule and the corresponding panel on Grafana dashboard.

## Multitenancy

There are the following approaches exist for alerting and recording rules across [multiple tenants](#):

- To run a separate `vmalert` instance per each tenant. The corresponding tenant must be specified in `-datasource.url` command-line flag according to [these docs](#). For example, `/path/to/vmalert -datasource.url=http://vmselect:8481/select/123/prometheus` would run alerts against `AccountID=123`. For recording rules the `-remoteWrite.url` command-line flag must contain the url for the specific tenant as well. For example, `-remoteWrite.url=http://vminsert:8480/insert/123/prometheus` would write recording rules to `AccountID=123`.
- To use the [multitenant endpoint](#) [②](#) of `vminsert` as the `-remoteWrite.url` and `vmselect` as the `-datasource.url`, add `extra_label` with tenant ID as an HTTP URL parameter for each group. For example, run `vmalert` using `-datasource.url=http://vmselect:8481/select/multitenant/prometheus -remoteWrite.url=http://vminsert:8480/insert/multitenant/prometheus`, along with the rule group:



```
groups:
- name: rules_for_tenant_456:789
  params:
    extra_label: [vm_account_id=456,vm_project_id=789]
  rules:
    # Rules for accountID=456, projectID=789
```

The multitenant endpoint in `vmselect` is less efficient than [specifying tenants in URL](#).

For security considerations, it is recommended restricting access to multitenant endpoints only to trusted sources, since untrusted source may break per-tenant data by writing unwanted samples or get access to data of arbitrary tenants.



- To specify `tenant` parameter per each alerting and recording group if [enterprise version of vmaalert](#) is used with `-clusterMode` command-line flag. For example:



```
groups:
- name: rules_for_tenant_123
  tenant: "123"
  rules:
    # Rules for accountID=123

- name: rules_for_tenant_456:789
  tenant: "456:789"
  rules:
    # Rules for accountID=456, projectID=789
```

The results of alerting and recording rules contain `vm_account_id` and `vm_project_id` labels if `-clusterMode` is enabled. These labels can be used during [templating](#), and help to identify to which account or project the triggered alert or produced recording belongs.

If `-clusterMode` is enabled, then `-datasource.url`, `-remoteRead.url` and `-remoteWrite.url` must contain only the hostname without tenant id. For example: `-datasource.url=http://vmselect:8481`. `vmaalert` automatically adds the specified tenant to urls per each recording rule in this case.

If `-clusterMode` is enabled and the `tenant` in a particular group is missing, then the tenant value is obtained from `-defaultTenant.prometheus` or `-defaultTenant.graphite` depending on the `type` of the group.

The enterprise version of `vmaalert` is available in `vmutils-*-enterprise.tar.gz` files at [release page](#) and in `*-enterprise` tags at [Docker Hub](#) and [Quay](#).

## Reading rules from object storage

[Enterprise version](#) of `vmaalert` may read alerting and recording rules from object storage:

- `./bin/vmaalert -rule=s3://bucket/dir/alert.rules` would read rules from the given path at S3 bucket
- `./bin/vmaalert -rule=gs://bucket/dir/alert.rules` would read rules from the given path at GCS bucket

S3 and GCS paths support only matching by prefix, e.g. `s3://bucket/dir/rule_` matches all files with prefix `rule_` in the folder `dir`.

The following [command-line flags](#) can be used for fine-tuning access to S3 and GCS:

- `-s3.credsFilePath` - path to file with GCS or S3 credentials. Credentials are loaded from default locations if not set.
- `-s3.configFilePath` - path to file with S3 configs. Configs are loaded from default location if not set.
- `-s3.configProfile` - profile name for S3 configs. If no set, the value of the environment variable will be loaded (`AWS_PROFILE` or `AWS_DEFAULT_PROFILE`).
- `-s3.customEndpoint` - custom S3 endpoint for use with S3-compatible storages (e.g. MinIO). S3 is used if not set.
- `-s3.forcePathStyle` - prefixing endpoint with bucket name when set false, true by default.

## Topology examples

The following sections are showing how `vmaalert` may be used and configured for different scenarios.

Please note, not all flags in examples are required:

- `-remoteWrite.url` and `-remoteRead.url` are optional and are needed only if you have recording rules or want to store [alerts state](#) on `vmaalert` restarts;
- `-notifier.url` is optional and is needed only if you have alerting rules.

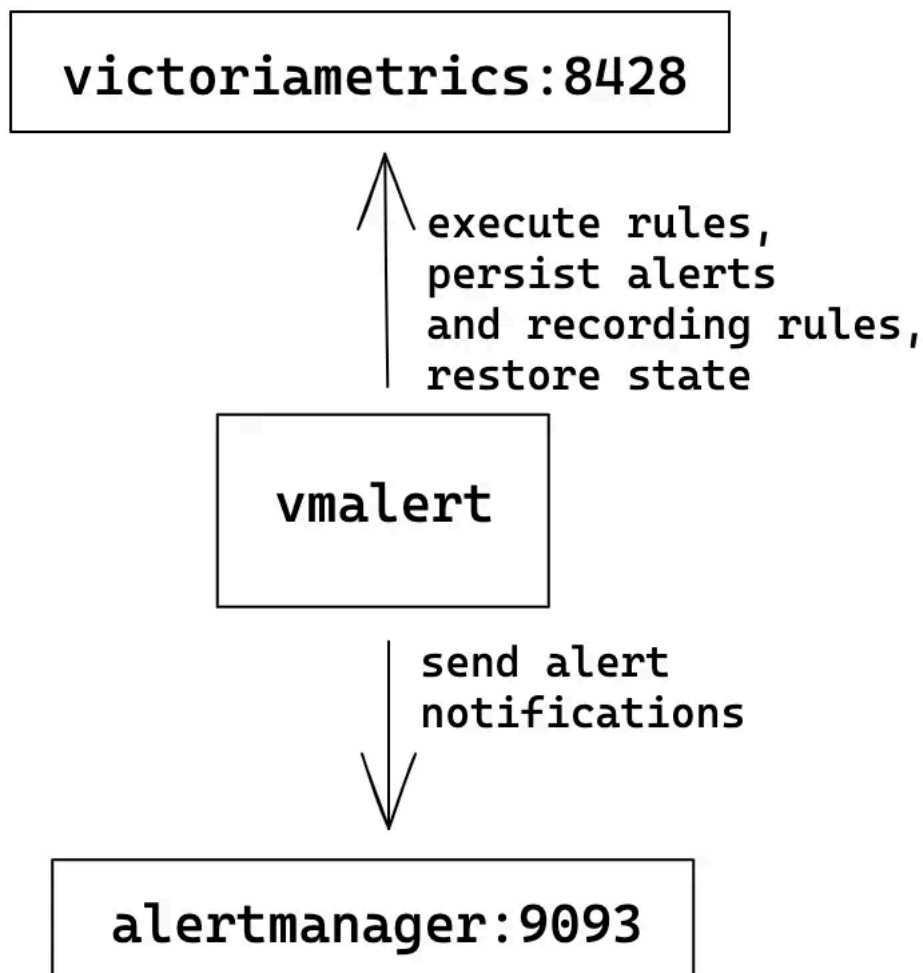
## Single-node VictoriaMetrics

The simplest configuration where one single-node VM server is used for rules execution, storing recording rules results and alerts state.

vmaalert configuration flags:



```
./bin/vmaalert -rule=rules.yml \           # Path to the file with rules configuration. Supports wildcard
-datasource.url=http://victoriametrics:8428 \ # VM-single addr for executing rules expressions
-remoteWrite.url=http://victoriametrics:8428 \ # VM-single addr to persist alerts state and recording rules results
-remoteRead.url=http://victoriametrics:8428 \ # VM-single addr for restoring alerts state after restart
-notifier.url=http://alertmanager:9093        # AlertManager addr to send alerts when they trigger
```



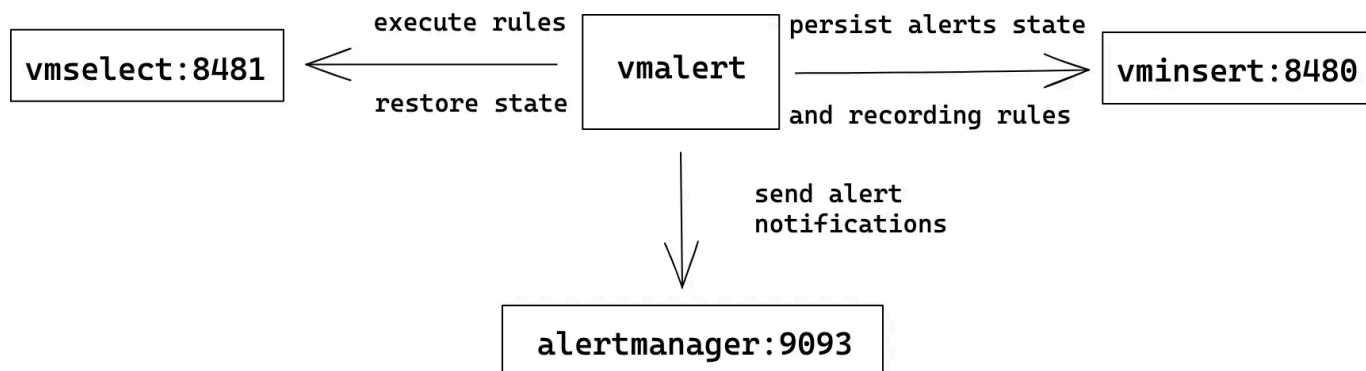
## Cluster VictoriaMetrics

In [cluster mode](#) VictoriaMetrics has separate components for writing and reading path: `vminsert` and `vmselect` components respectively. `vmselect` is used for executing rules expressions and `vminsert` is used to persist recording rules results and alerts state. Cluster mode could have multiple `vminsert` and `vmselect` components.

vmaalert configuration flags:



```
./bin/vmaalert -rule=rules.yml \           # Path to the file with rules configuration. Supports wildcard
-datasource.url=http://vmselect:8481/select/0/prometheus # vmselect addr for executing rules expressions
-remoteWrite.url=http://vminsert:8480/insert/0/prometheus # vminsert addr to persist alerts state and recording rules results
-remoteRead.url=http://vmselect:8481/select/0/prometheus # vmselect addr for restoring alerts state after restart
-notifier.url=http://alertmanager:9093        # AlertManager addr to send alerts when they trigger
```



In case when you want to spread the load on these components - add balancers before them and configure `vmaalert` with balancer addresses. Please, see more about VM's cluster architecture [here](#).

## HA vmaalert

For High Availability(HA) user can run multiple identically configured `vmaalert` instances. It means all of them will execute the same rules, write state and results to the same destinations, and send alert notifications to multiple configured Alertmanagers.

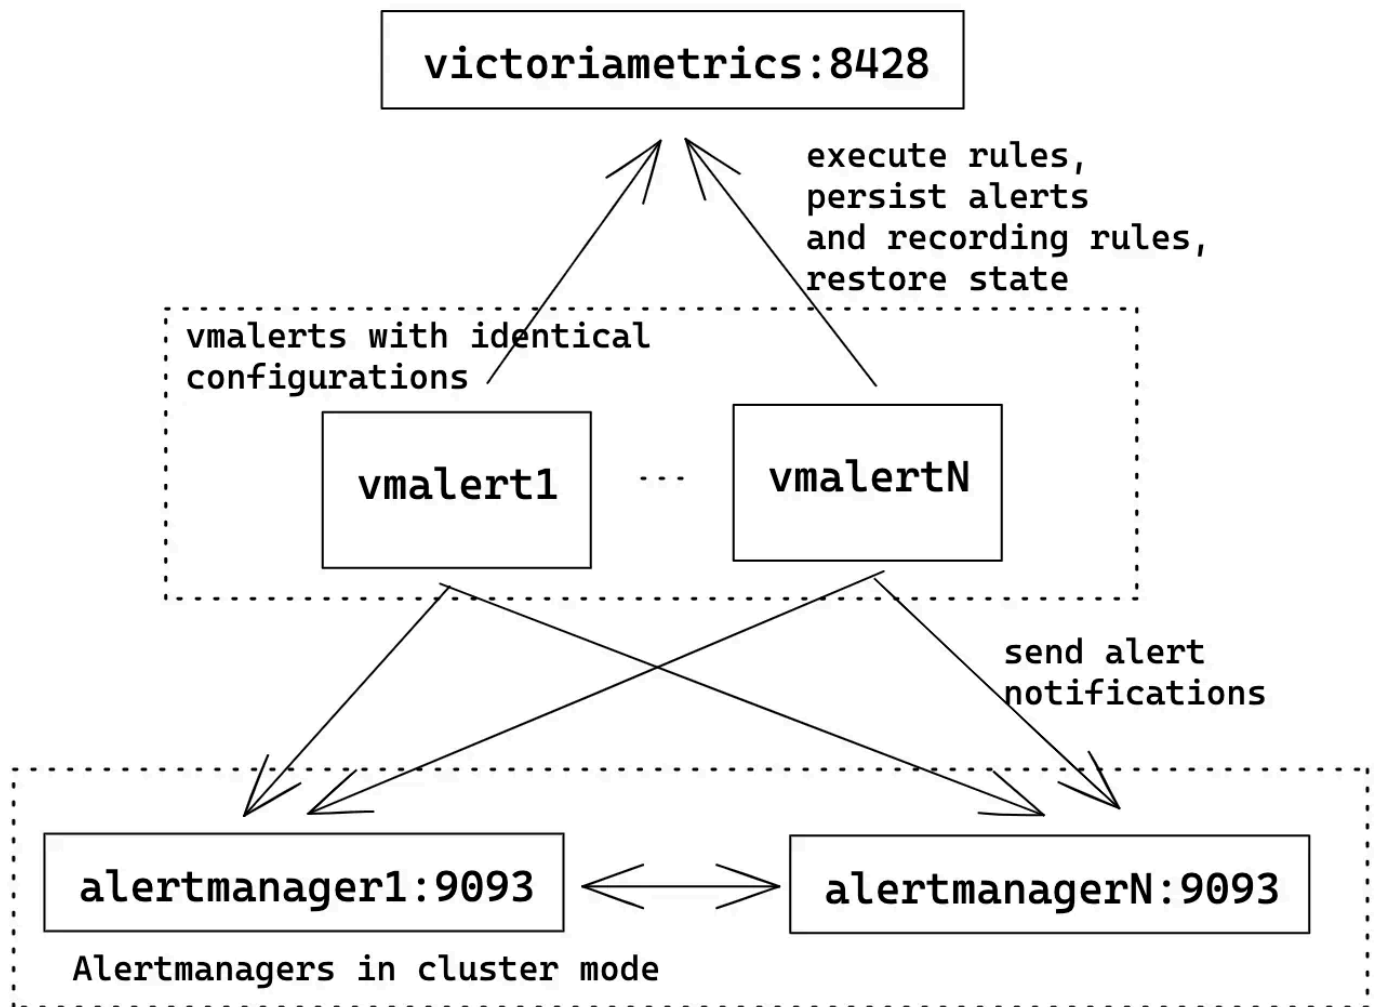
`vmaalert` configuration flags:



```

./bin/vmaalert -rule=rules.yml \
  -datasource.url=http://victoriametrics:8428 \ # Path to the file with rules configuration. Supports wildcard
  -remoteWrite.url=http://victoriametrics:8428 \ # VM-single addr for executing rules expressions
  -remoteRead.url=http://victoriametrics:8428 \ # VM-single addr to persist alerts state and recording rules results
  -notifier.url=http://alertmanager1:9093 \ # VM-single addr for restoring alerts state after restart
  -notifier.url=http://alertmanagerN:9093 \ # Multiple AlertManager addresses to send alerts when they trigger
                                           # The same alert will be sent to all configured notifiers
  
```

## VictoriaMetrics with deduplication enabled



To avoid recording rules results and alerts state duplication in VictoriaMetrics server don't forget to configure [deduplication](#). Multiple equally configured vmaalerts should evaluate rules at the same timestamps, so it is recommended to set `-dedup.minScrapeInterval` as equal to vmaalert's `-evaluationInterval`.

If you have multiple different `interval` params for distinct rule groups, then set `-dedup.minScrapeInterval` to the biggest `interval` value, or value which will be a multiple for all `interval` values. For example, if you have two groups with `interval: 10s` and `interval: 15s`, then set `-dedup.minScrapeInterval=30s`. This would consistently keep only a single data point on 30s time interval for all rules. However, try to avoid having inconsistent `interval` values.

It is not recommended having `-dedup.minScrapeInterval` smaller than `-evaluationInterval`, as it may produce results with inconsistent intervals between data points.

Alertmanager will automatically deduplicate alerts with identical labels, so ensure that all vmaalert s are having identical config.

Don't forget to configure [cluster mode](#) for Alertmanagers for better reliability. List all Alertmanager URLs in vmaalert `-notifier.url` to ensure [high availability](#).

This example uses single-node VM server for the sake of simplicity. Check how to replace it with [cluster VictoriaMetrics](#) if needed.

### Downsampling and aggregation via vmaalert

Please note, [stream aggregation](#) might be more efficient for cases when downsampling or aggregation need to be applied **before data gets into the TSDB**.

vmaalert can't modify existing data. But it can run arbitrary PromQL/MetricsQL queries via [recording rules](#) and backfill results to the configured `-remoteWrite.url`. This ability allows to aggregate data. For example, the following rule will calculate the average value for metric `http_requests` on the 5m interval:



```
- record: http_requests:avg5m
  expr: avg_over_time(http_requests[5m])
```

Every time this rule will be evaluated, `vmaalert` will backfill its results as a new time series `http_requests:avg5m` to the configured `-remoteWrite.url`.

`vmaalert` executes rules with specified interval (configured via flag `-evaluationInterval` or as `group's interval` param). The interval helps to control “resolution” of the produced series. This ability allows to downsample data. For example, the following config will execute the rule only once every `5m`:



```
groups:
- name: my_group
  interval: 5m
  rules:
- record: http_requests:avg5m
  expr: avg_over_time(http_requests[5m])
```

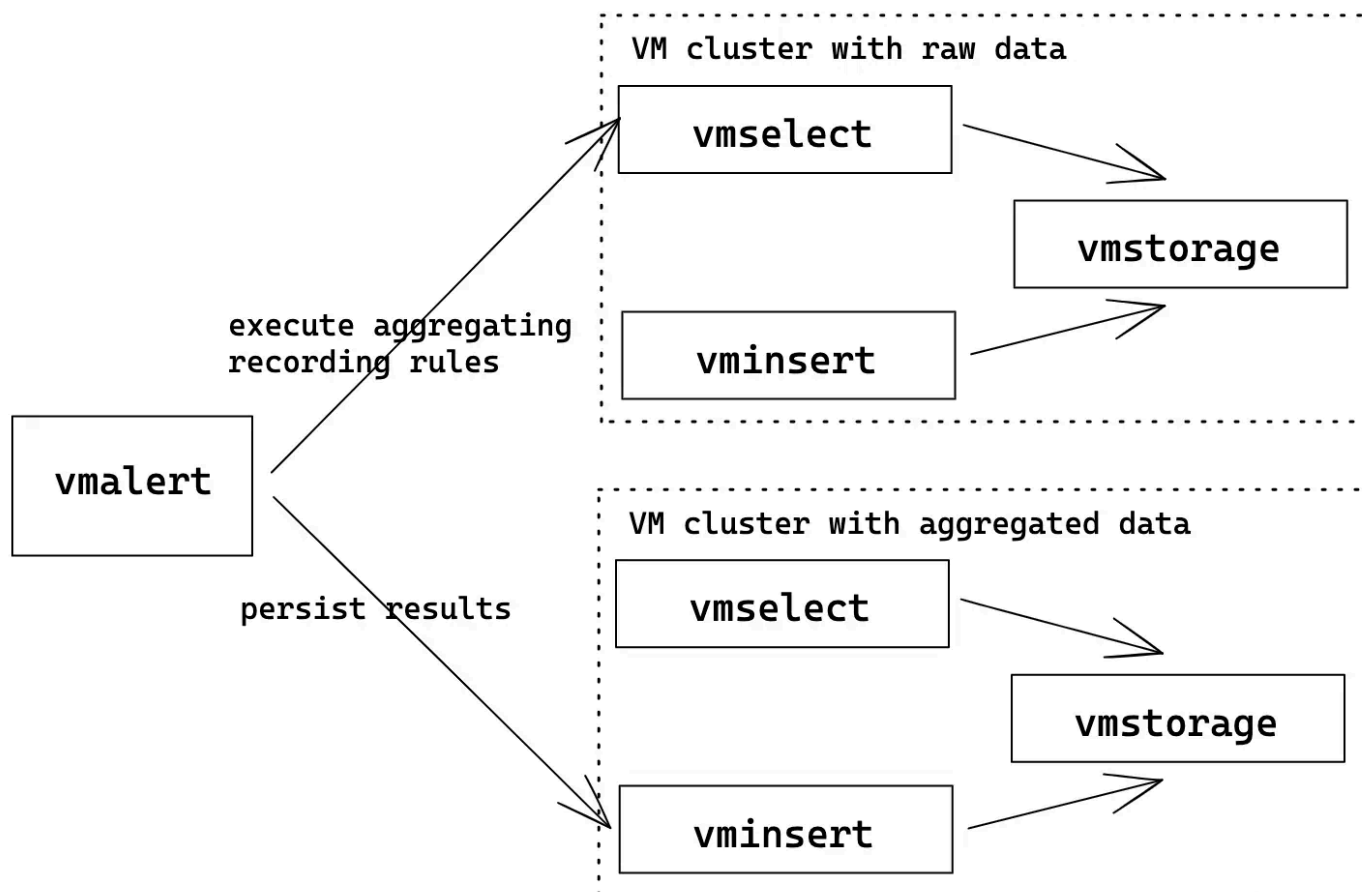
Ability of `vmaalert` to be configured with different `-datasource.url` and `-remoteWrite.url` command-line flags allows reading data from one data source and backfilling results to another. This helps to build a system for aggregating and downsampling the data.

The following example shows how to build a topology where `vmaalert` will process data from one cluster and write results into another. Such clusters may be called as “hot” (low retention, high-speed disks, used for operative monitoring) and “cold” (long term retention, slower/cheaper disks, low resolution data). With help of `vmaalert`, user can setup recording rules to process raw data from “hot” cluster (by applying additional transformations or reducing resolution) and push results to “cold” cluster.

`vmaalert` configuration flags:



```
./bin/vmaalert -rule=downsampling-rules.yml \
-datasource.url=http://raw-cluster-vmselect:8481/select/0/prometheus # Path to the file with rules configuration. Supports wildcards
-remoteWrite.url=http://aggregated-cluster-vminsert:8480/insert/0/prometheus # vmselect addr for executing recording rules expressions
# vminsert addr to persist recording rules results
```




Please note, [replay](#) feature may be used for transforming historical data.

Flags `-remoteRead.url` and `-notifier.url` are omitted since we assume only recording rules are used.

See also [stream aggregation](#) and [downsampling](#).

### Multiple remote writes

For persisting recording or alerting rule results `vmaalert` requires `-remoteWrite.url` to be set. But this flag supports only one destination. To persist rule results to multiple destinations we recommend using [vmagent](#) as fan-out proxy:



In this topology, `vmaalert` is configured to persist rule results to `vmagent`. And `vmagent` is configured to fan-out received data to two or more destinations. Using `vmagent` as a proxy provides additional benefits such as [data persisting when storage is unreachable](#), or time series modification via [relabeling](#).

### Web

`vmaalert` runs a web-server ( `-httpListenAddr` ) for serving metrics and alerts endpoints:

- `http://<vmaalert-addr>` - UI;
- `http://<vmaalert-addr>/api/v1/rules` - list of all loaded groups and rules. Supports additional [filtering](#);
- `http://<vmaalert-addr>/api/v1/alerts` - list of all active alerts;

- `http://<vmaalert-addr>/vmaalert/api/v1/alert?group_id=<group_id>&alert_id=<alert_id>` - get alert status in JSON format. Used as alert source in AlertManager.
- `http://<vmaalert-addr>/vmaalert/alert?group_id=<group_id>&alert_id=<alert_id>` - get alert status in web UI.
- `http://<vmaalert-addr>/vmaalert/rule?group_id=<group_id>&rule_id=<rule_id>` - get rule status in web UI.
- `http://<vmaalert-addr>/vmaalert/api/v1/rule?group_id=<group_id>&alert_id=<alert_id>` - get rule status in JSON format.
- `http://<vmaalert-addr>/metrics` - application metrics.
- `http://<vmaalert-addr>/-/reload` - hot configuration reload.

vmaalert web UI can be accessed from [single-node version of VictoriaMetrics](#) and from [cluster version of VictoriaMetrics](#). This may be used for better integration with Grafana unified alerting system. See the following docs for details:

- [How to query vmaalert from single-node VictoriaMetrics](#)
- [How to query vmaalert from VictoriaMetrics cluster](#)

## Graphite

vmaalert sends requests to `<-datasource.url>/render?format=json` during evaluation of alerting and recording rules if the corresponding group or rule contains `type: "graphite"` config option. It is expected that the `<-datasource.url>/render` implements [Graphite Render API](#) for `format=json`. When using vmaalert with both `graphite` and `prometheus` rules configured against cluster version of VM do not forget to set `-datasource.appendTypePrefix` flag to `true`, so vmaalert can adjust URL prefix automatically based on the query type.

## VictoriaLogs

vmaalert supports [VictoriaLogs](#) as a datasource for writing alerting and recording rules using [LogsQL](#). See [this doc](#) for details.

## Rules backfilling

vmaalert supports alerting and recording rules backfilling (aka `replay`). In replay mode vmaalert can read the same rules configuration as normal, evaluate them on the given time range and backfill results via remote write to the configured storage. vmaalert supports any PromQL/MetricsQL compatible data source for backfilling.

Please note, that response caching may lead to unexpected results during and after backfilling process. In order to avoid this you need to reset cache contents or disable caching when using backfilling as described in [backfilling docs](#).

See a blogpost about [Rules backfilling via vmaalert](#).

## How it works

In `replay` mode vmaalert works as a cli-tool and exits immediately after work is done. To run vmaalert in `replay` mode:



```
./bin/vmaalert -rule=path/to/your.rules \      # path to files with rules you usually use with vmaalert
  -datasource.url=http://localhost:8428 \      # Prometheus HTTP API compatible datasource
  -remoteWrite.url=http://localhost:8428 \      # remote write compatible storage to persist results
  -replay.timeFrom=2021-05-11T07:21:43Z \      # to start replay from
  -replay.timeTo=2021-05-29T18:40:43Z          # to finish replay by, is optional
```

The output of the command will look like the following:



```
Replay mode:
from: 2021-05-11 07:21:43 +0000 UTC # set by -replay.timeFrom
to: 2021-05-29 18:40:43 +0000 UTC # set by -replay.timeTo
```

```

max data points per request: 1000      # set by -replay.maxDatapointsPerQuery

Group "ReplayGroup"
interval:      1m0s
requests to make:      27
max range per request: 16h40m0s
> Rule "type:vm_cache_entries:rate5m" (ID: 1792509946081842725)
27 / 27 [-----] 100.00% ? p/s
> Rule "go_cgo_calls_count:rate5m" (ID: 17958425467471411582)
27 / 27 [-----] 100.00% ? p/s

Group "vmsingleReplay"
interval:      30s
requests to make:      54
max range per request: 8h20m0s
> Rule "RequestErrorsToAPI" (ID: 17645863024999990222)
54 / 54 [-----] 100.00% ? p/s
> Rule "TooManyLogs" (ID: 9042195394653477652)
54 / 54 [-----] 100.00% ? p/s
2021-06-07T09:59:12.098Z      info      app/vmalert/replay.go:68      replay finished! Imported 511734 samples

```

In `replay` mode all groups are executed sequentially one-by-one. Rules within the group are executed sequentially as well ( `concurrency` setting is ignored). `vmalert` sends rule's expression to `/query_range` endpoint of the configured `-datasource.url`. Returned data is then processed according to the rule type and backfilled to `-remoteWrite.url` via [remote Write protocol](#). `vmalert` respects `evaluationInterval` value set by flag or per-group during the replay. `vmalert` automatically disables caching on VictoriaMetrics side by sending `nocache=1` param. It allows to prevent cache pollution and unwanted time range boundaries adjustment during backfilling.

## Recording rules

The result of recording rules `replay` should match with results of normal rules evaluation.

## Alerting rules

The result of alerting rules `replay` is time series reflecting [alert's state](#). To see if `replayed` alert has fired in the past use the following PromQL/MetricsQL expression:

```
ALERTS{alertname="your_alertname", alertstate="firing"}
```



Execute the query against storage which was used for `-remoteWrite.url` during the `replay`.

## Additional configuration

There are following non-required `replay` flags:

- `-replay.maxDatapointsPerQuery` - the max number of data points expected to receive in one request. In two words, it affects the max time range for every `/query_range` request. The higher the value, the fewer requests will be issued during `replay`.
- `-replay.ruleRetryAttempts` - when `datasource` fails to respond `vmalert` will make this number of retries per rule before giving up.
- `-replay.rulesDelay` - delay between sequential rules execution. Important in cases if there are chaining (rules which depend on each other) rules. It is expected, that remote storage will be able to persist previously accepted data during the delay, so data will be available for the subsequent queries. Keep it equal or bigger than `-remoteWrite.flushInterval`.
- `-replay.disableProgressBar` - whether to disable progress bar which shows progress work. Progress bar may generate a lot of log records, which is not formatted as standard VictoriaMetrics logger. It could break logs parsing by external system and generate additional load on it.

See full description for these flags in `./vmalert -help`.

## Limitations

- Graphite engine isn't supported yet;



- query template function is disabled for performance reasons (might be changed in future);
- limit\_group's param has no effect during replay (might be changed in future);
- keep\_firing\_for alerting rule param has no effect during replay (might be changed in future).

## Unit Testing for Rules

You can use `vmaAlert-tool` to test your alerting and recording rules like [promtool does](#). See more details [here](#).

## Monitoring

`vmaAlert` exports various metrics in Prometheus exposition format at `http://vmaAlert-host:8880/metrics` page. See the [list of recommended alerting rules](#) to track the `vmaAlert` health. We recommend setting up regular scraping of this page either through [vmagent](#) or by Prometheus-compatible scraper, so that the exported metrics may be analyzed later.

If you use Google Cloud Managed Prometheus for scraping metrics from VictoriaMetrics components, then pass `-metrics.exposeMetadata` command-line to them, so they add `TYPE` and `HELP` comments per each exposed metric at `/metrics` page. See [these docs](#) for details.

Use the official [Grafana dashboard](#) for `vmaAlert` overview. Graphs on this dashboard contain useful hints - hover the `i` icon in the top left corner of each graph in order to read it. If you have suggestions for improvements or have found a bug - please open an issue on GitHub or add a review to the dashboard.

## Troubleshooting

### Rule state

`vmaAlert` stores last `-rule.updateEntriesLimit` (or `update_entries_limit` [per-rule config](#)) evaluation updates for each rule. The updates are available in `vmaAlert` [web UI](#):

1. Go to the `Groups` tab
2. Find the corresponding Group and rule within it
3. Click on `Details` link next to rule's name and check the `Last N updates` section:

Rule: code:requests:slo ok

Expr	0.95
Labels	<span>cluster=east-1</span> <span>recording=true</span> <span>replica=a</span>
Group	<a href="#">13681179853670628447</a>

### Last 13 updates:

Updated at	Samples	Duration	Executed at	cURL
2022-09-15T15:02:48+02:00	1	0.001s	2022-09-15T15:02:48+02:00	curl -X POST -H 'Content-Type: application/json' -H 'Myheader: foo' 'http://localhost:8428/api/v1/query?' <span>...</span>
2022-09-15T15:02:43+02:00	1	0.000s	2022-09-15T15:02:43+02:00	curl -X POST -H 'Content-Type: application/json' -H 'Myheader: foo' 'http://localhost:8428/api/v1/query?' <span>...</span>
2022-09-15T15:02:38+02:00	1	0.001s	2022-09-15T15:02:38+02:00	curl -X POST -H 'Content-Type: application/json' -H 'Myheader: foo' 'http://localhost:8428/api/v1/query?' <span>...</span>

Rows in the section represent ordered rule evaluations and their results.

Every state has the following attributes:

1. `Updated at` - the real time when this rule was executed by `vmaAlert`.
2. `Executed at` - `time` param that was sent within the evaluation request to datasource.

3. `Samples` - the amount of data samples returned during at this evaluation. Recording rule that has 0 samples produces no results. Alerting rule that has 0 samples means `expr` condition is false and rule is in inactive state.
4. `Series fetched` - the amount of series scanned during query evaluation. See [never-firing alerts](#).
5. `Duration` - the time it took to evaluate the rule. If `Duration` is close or bigger to evaluation interval, then this rule can skip evaluations. See how to deal with [slow queries](#).
6. `cURL` - contains an example of HTTP request sent by vmaalert to the `-datasource.url` during evaluation, including all extra headers and query params. The command can be used for debugging purposes to see what vmaalert receives in response from datasource. *Sensitive info is stripped from the `cURL` examples - see [security](#) section for more details.*

If a specific state shows that there were **no samples returned** and executed **cURL command returns some data**, then it is likely there was no data in datasource on the moment when rule was evaluated. See about [data delay](#).

vmaalert exposes `vmaalert_recording_rules_last_evaluation_samples` for recording rules to represent the amount of samples returned during evaluations. The following alerting rule can be used to detect those recording rules that produce no data:



```
- alert: RecordingRulesNoData
  expr: vmaalert_recording_rules_last_evaluation_samples < 1
  annotations:
    summary: "Recording rule {{ $labels.recording }} ({{ $labels.group }}) produces no data"
```

See more about alerting rules in [Monitoring](#).

## Alert state

Sometimes, it is not clear why some specific alert fired or didn't fire. It is very important to remember, that alerts with `for: 0` fire immediately when their expression becomes true. And alerts with `for: > 0` will fire only after **multiple consecutive evaluations**, where each evaluation of the expression must be positive.

If at least one evaluation returns no data, then alert's state resets to the initial state.


### Note

Please note, the alert state is tracked individually for each returned time series during evaluation. If 1st evaluation returned series A and B, and 2nd evaluation returned only B - then alert will remain active only for B.


If `-remoteWrite.url` command-line flag is configured, vmaalert will [persist alert's state](#) in form of time series `ALERTS` and `ALERTS_FOR_STATE` to the specified destination. Such time series can be then queried via [vmui](#) or Grafana to track how alerts state changed in time.

## Data delay

Data delay is one of the most common issues with rules execution. vmaalert executes configured rules within certain intervals at specific timestamps. It expects that the data is already present in configured `-datasource.url` at the moment of time when rule is executed:



Usually, troubles begin when data in `-datasource.url` is delayed or absent. In such cases, evaluations may get empty response from the datasource and produce empty recording rules or reset alerts state:



Try the following recommendations to reduce the chance of hitting the data delay issue:

- Always configure group's `-evaluationInterval` to be bigger or at least equal to [time series resolution](#) ;
- Ensure that `[duration]` value is at least twice bigger than [time series resolution](#) . For example, if expression is `rate(my_metric[2m]) > 0` then ensure that `my_metric` resolution is at least 1m or better 30s .
- Extend `[duration]` in expr to help tolerate the delay. For example, `max_over_time(errors_total[10m]) > 0` will be active even if there is no data in datasource for last 9m .
- If [time series resolution](#) in datasource is inconsistent or `>=5min` - try changing vmaalerts `-datasource.queryStep` command-line flag to specify how far search query can look back for the recent datapoint. The recommendation is to have the step at least two times bigger than the resolution.

#### Note

Please note, data delay is inevitable in distributed systems. And it is better to account for it instead of ignoring.

By default, recently written samples to VictoriaMetrics [aren't visible for queries](#) for up to 30s (see `-search.latencyOffset` command-line flag at `vmselect` or VictoriaMetrics single-node). Such delay is needed to eliminate risk of incomplete data on the moment of querying, due to chance that metrics collectors won't be able to deliver that data in time. To compensate the latency in timestamps for produced evaluation results, `-rule.evalDelay` is also set to 30s by default. If you expect data to be delayed for longer intervals (it gets buffered, queued, or just network is slow sometimes), or you changed default value of `-search.latencyOffset` - consider increasing the `-rule.evalDelay` value accordingly.

## Debug mode

vmaalert allows configuring more detailed logging for specific alerting or recording rule starting from [v1.91](#) . Or for all rules within the [group](#) [v1.91](#) . Just set `debug: true` in configuration and vmaalert will start printing additional log messages:



```
2022-09-15T13:35:41.155Z DEBUG alerting rule "TestGroup":"Conns" (2601299393013563564) at 2022-09-15T15:35:41+02:00: query returned 0 samples (
2022-09-15T13:35:56.149Z DEBUG datasource request: executing POST request with params "denyPartialResponse=true&query=sum%28vm_tcpllistener_conn:
2022-09-15T13:35:56.178Z DEBUG alerting rule "TestGroup":"Conns" (2601299393013563564) at 2022-09-15T15:35:56+02:00: query returned 1 samples (
2022-09-15T13:35:56.178Z DEBUG datasource request: executing POST request with params "denyPartialResponse=true&query=sum%28vm_tcpllistener_conn:
2022-09-15T13:35:56.179Z DEBUG alerting rule "TestGroup":"Conns" (2601299393013563564) at 2022-09-15T15:35:56+02:00: alert 10705778000901301787
...
2022-09-15T13:36:56.153Z DEBUG alerting rule "TestGroup":"Conns" (2601299393013563564) at 2022-09-15T15:36:56+02:00: alert 10705778000901301787
```

Sensitive info is stripped from the `curl` examples - see [security](#) section for more details.

## Never-firing alerts

vmaalert can detect [v1.91](#) if alert's expression doesn't match any time series in runtime starting from [v1.91](#) . This problem usually happens when alerting expression selects time series which aren't present in the datasource (i.e. wrong `job` label) or there is a typo in the series selector (i.e. `env=prod` ). Such alerting rules will be marked with special icon in vmaalerts UI and exposed via `vmaalert_alerting_rules_last_evaluation_series_fetched` metric. The metric value will show how many time series were matched before the filtering by rule's expression. If metric value is `-1` , then this feature is not supported by the datasource (old versions of VictoriaMetrics). The following expression can be used to detect rules matching no series:



```
max(vmaalert_alerting_rules_last_evaluation_series_fetched) by(group, alertname) == 0
```

See more details [here](#) and read [Never-firing alerts](#) blogpost.

## Series with the same labelset

vmaalert can produce the following error message:



```
result contains metrics with the same labelset during evaluation
```

The error means there is a collision between [time series](#) during evaluation.

For example, a rule with `expr: {__name__=~"vmaalert_alerts_.*"} > 0` returns two distinct time series in response:



```
{__name__="vmaalert_alerts_pending",job="vmaalert",alertname="HostContextSwitching"} 12
{__name__="vmaalert_alerts_firing",job="vmaalert",alertname="HostContextSwitching"} 0
```

As label `__name__` will be dropped during evaluation, leads to duplicated time series. To fix this, one could use function like [label\\_replace](#) to preserve the distinct labelset.

## mTLS protection

By default `vmaalert` accepts http requests at `8880` port (this port can be changed via `-httpListenAddr` command-line flags), since it is expected it runs in an isolated trusted network. [Enterprise version of vmagent](#) supports the ability to accept [mTLS](#) requests at this port, by specifying `-tls` and `-mtls` command-line flags. For example, the following command runs `vmaalert`, which accepts only mTLS requests at port `8880`:



```
./vmaalert -tls -mtls -remoteWrite.url=...
```

By default system-wide [TLS Root CA](#) is used for verifying client certificates if `-mtls` command-line flag is specified. It is possible to specify custom TLS Root CA via `-mtlsCAFile` command-line flag.

## Security

See general recommendations regarding security [here](#).

vmaalert [web UI](#) exposes configuration details such as list of [Groups](#), active alerts, [alerts state](#), [notifiers](#). Notifier addresses (sanitized) are attached as labels to metrics `vmaalert_alerts_sent_.*` on `http://<vmaalert>/metrics` page. Consider limiting user's access to the web UI or `/metrics` page if this information is sensitive.

[Alerts state](#) page or [debug mode](#) could emit additional information about configured datasource URL, GET params and headers. Sensitive information such as passwords or auth tokens is stripped by default. To disable stripping of such info pass `-datasource.showURL` cmd-line flag to vmaalert.

See also [mTLS protection docs](#).

## Profiling

vmaalert provides handlers for collecting the following [Go profiles](#):

- Memory profile. It can be collected with the following command (replace `0.0.0.0` with hostname if needed):



```
curl http://0.0.0.0:8880/debug/pprof/heap > mem.pprof
```

- CPU profile. It can be collected with the following command (replace 0.0.0.0 with hostname if needed):



```
curl http://0.0.0.0:8880/debug/pprof/profile > cpu.pprof
```

The command for collecting CPU profile waits for 30 seconds before returning.

The collected profiles may be analyzed with [go tool pprof](#). It is safe sharing the collected profiles from security point of view, since they do not contain sensitive information.

## Configuration

### Flags

Pass `-help` to `vmaalert` in order to see the full list of supported command-line flags with their descriptions.

The shortlist of configuration flags is the following:



```
-blockcache.missesBeforeCaching int
    The number of cache misses before putting the block into cache. Higher values may reduce indexdb/dataBlocks cache size at the cost of higher
-clusterMode
    If clusterMode is enabled, then vmaalert automatically adds the tenant specified in config groups to -datasource.url, -remoteWrite.url and -r
-configCheckInterval duration
    Interval for checking for changes in '-rule' or '-notifier.config' files. By default, the checking is disabled. Send SIGHUP signal in order
-datasource.appendTypePrefix
    Whether to add type prefix to -datasource.url based on the query type. Set to true if sending different query types to the vmselect URL.
-datasource.basicAuth.password string
    Optional basic auth password for -datasource.url
-datasource.basicAuth.passwordFile string
    Optional path to basic auth password to use for -datasource.url
-datasource.basicAuth.username string
    Optional basic auth username for -datasource.url
-datasource.bearerToken string
    Optional bearer auth token to use for -datasource.url.
-datasource.bearerTokenFile string
    Optional path to bearer token file to use for -datasource.url.
-datasource.disableKeepAlive
    Whether to disable long-lived connections to the datasource. If true, disables HTTP keep-alive and will only use the connection to the servi
-datasource.disableStepParam
    Whether to disable adding 'step' param in instant queries to the configured -datasource.url and -remoteRead.url. Only valid for prometheus
-datasource.headers string
    Optional HTTP extraHeaders to send with each request to the corresponding -datasource.url. For example, -datasource.headers='My-Auth:foobar
-datasource.idleConnTimeout duration
    Defines a duration for idle (keep-alive connections) to exist. Consider setting this value less than "-http.idleConnTimeout". It must preven
-datasource.maxIdleConnections int
    Defines the number of idle (keep-alive connections) to each configured datasource. Consider setting this value equal to the value: groups_t
-datasource.oauth2.clientID string
    Optional OAuth2 clientID to use for -datasource.url
-datasource.oauth2.clientSecret string
    Optional OAuth2 clientSecret to use for -datasource.url
-datasource.oauth2.clientSecretFile string
    Optional OAuth2 clientSecretFile to use for -datasource.url
-datasource.oauth2.endpointParams string
    Optional OAuth2 endpoint parameters to use for -datasource.url . The endpoint parameters must be set in JSON format: {"param1":"value1",...
-datasource.oauth2.scopes string
    Optional OAuth2 scopes to use for -datasource.url. Scopes must be delimited by ';'
-datasource.oauth2.tokenUrl string
    Optional OAuth2 tokenURL to use for -datasource.url
-datasource.queryStep duration
    How far a value can fallback to when evaluating queries to the configured -datasource.url and -remoteRead.url. Only valid for prometheus da
-datasource.roundDigits int
```

Adds "round\_digits" GET param to datasource requests which limits the number of digits after the decimal point in response values. Only valid for Prometheus datasource.

Whether to avoid stripping sensitive information such as auth headers or passwords from URLs in log messages or UI and exported metrics. It is only valid for Prometheus datasource.

-datasource.tlsCAFile string  
Optional path to TLS CA file to use for verifying connections to -datasource.url. By default, system CA is used

-datasource.tlsCertFile string  
Optional path to client-side TLS certificate file to use when connecting to -datasource.url

-datasource.tlsInsecureSkipVerify  
Whether to skip tls verification when connecting to -datasource.url

-datasource.tlsKeyFile string  
Optional path to client-side TLS certificate key to use when connecting to -datasource.url

-datasource.tlsServerName string  
Optional TLS server name to use for connections to -datasource.url. By default, the server name from -datasource.url is used

-datasource.url string  
Datasource compatible with Prometheus HTTP API. It can be single node VictoriaMetrics or vmselect endpoint. Required parameter. Supports additional parameters for Prometheus datasource.

-defaultTenant.graphite string  
Default tenant for Graphite alerting groups. See <https://docs.victoriametrics.com/victoriametrics/vmalert/#multitenancy>. This flag is available only for Prometheus datasource.

-defaultTenant.prometheus string  
Default tenant for Prometheus alerting groups. See <https://docs.victoriametrics.com/victoriametrics/vmalert/#multitenancy>. This flag is available only for Prometheus datasource.

-disableAlertgroupLabel  
Whether to disable adding group's Name as label to generated alerts and time series.

-dryRun  
Whether to check only config files without running vmalert. The rules file are validated. The -rule flag must be specified.

-enableTCP6  
Whether to enable IPv6 for listening and dialing. By default, only IPv4 TCP and UDP are used

-envflag.enable  
Whether to enable reading flags from environment variables in addition to the command line. Command line flag values have priority over environment variables.

-envflag.prefix string  
Prefix for environment variables if -envflag.enable is set

-eula  
Deprecated, please use -license or -licenseFile flags instead. By specifying this flag, you confirm that you have an enterprise license and agree to the terms of the license.

-evaluationInterval duration  
How often to evaluate the rules (default 1m0s)

-external.alert.source string  
External Alert Source allows to override the Source link for alerts sent to AlertManager for cases where you want to build a custom link to the source of the alert.

-external.label exported\_label  
Optional label in the form 'Name=value' to add to all generated recording rules and alerts. In case of conflicts, original labels are kept. Supports an 'array' of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

-external.url string  
External URL is used as alert's source for sent alerts to the notifier. By default, hostname is used as address.

-filestream.disableFadvise  
Whether to disable fadvise() syscall when reading large data files. The fadvise() syscall prevents from eviction of recently accessed data from memory.

-flagsAuthKey value  
Auth key for /flags endpoint. It must be passed via authKey query arg. It overrides -httpAuth.\*  
Flag value can be read from the given file when using -flagsAuthKey=file:///abs/path/to/file or -flagsAuthKey=file:///relative/path/to/file

-fs.disableMmap  
Whether to use pread() instead of mmap() for reading data files. By default, mmap() is used for 64-bit arches and pread() is used for 32-bit arches.

-http.connTimeout duration  
Incoming connections to -httpListenAddr are closed after the configured timeout. This may help evenly spreading load among a cluster of servers.

-http.disableCORS  
Disable CORS for all origins (\*)

-http.disableResponseCompression  
Disable compression of HTTP responses to save CPU resources. By default, compression is enabled to save network bandwidth

-http.header.csp string  
Value for 'Content-Security-Policy' header, recommended: "default-src 'self'"

-http.header.frameOptions string  
Value for 'X-Frame-Options' header

-http.header.hsts string  
Value for 'Strict-Transport-Security' header, recommended: 'max-age=31536000; includeSubDomains'

-http.idleConnTimeout duration  
Timeout for incoming idle http connections (default 1m0s)

-http.maxGracefulShutdownDuration duration  
The maximum duration for a graceful shutdown of the HTTP server. A highly loaded server may require increased value for a graceful shutdown

-http.pathPrefix string  
An optional prefix to add to all the paths handled by http server. For example, if '-http.pathPrefix=/foo/bar' is set, then all the http requests will be handled by paths starting with /foo/bar

-http.shutdownDelay duration  
Optional delay before http server shutdown. During this delay, the server returns non-OK responses from /health page, so load balancers can gracefully shutdown the server

-httpAuth.password value  
Password for HTTP server's Basic Auth. The authentication is disabled if -httpAuth.username is empty  
Flag value can be read from the given file when using -httpAuth.password=file:///abs/path/to/file or -httpAuth.password=file:///relative/path/to/file

-httpAuth.username string  
Username for HTTP server's Basic Auth. The authentication is disabled if empty. See also -httpAuth.password

-httpListenAddr array  
Address to listen for incoming http requests. See also -tls and -httpListenAddr.useProxyProtocol  
Supports an array of values separated by comma or specified via multiple flags.  
Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

-httpListenAddr.useProxyProtocol array  
Whether to use proxy protocol for incoming http requests. See also -httpListenAddr

Whether to use proxy protocol for connections accepted at the corresponding `-httpListenAddr` . See <https://www.haproxy.org/download/1.8/doc/>  
 Supports array of values separated by comma or specified via multiple flags.  
 Empty values are set to false.

`-internStringCacheExpireDuration` duration  
 The expiry duration for caches for interned strings. See [https://en.wikipedia.org/wiki/String\\_interning](https://en.wikipedia.org/wiki/String_interning) . See also `-internStringMaxLen` and `-internStringDisableCache`

`-internStringDisableCache`  
 Whether to disable caches for interned strings. This may reduce memory usage at the cost of higher CPU usage. See <https://en.wikipedia.org/>

`-internStringMaxLen` int  
 The maximum length for strings to intern. A lower limit may save memory at the cost of higher CPU usage. See <https://en.wikipedia.org/wiki/>

`-license` string  
 License key for VictoriaMetrics Enterprise. See <https://victoriametrics.com/products/enterprise/> . Trial Enterprise license can be obtained

`-license.forceOffline`  
 Whether to enable offline verification for VictoriaMetrics Enterprise license key, which has been passed either via `-license` or via `-licenseFile`

`-licenseFile` string  
 Path to file with license key for VictoriaMetrics Enterprise. See <https://victoriametrics.com/products/enterprise/> . Trial Enterprise licen

`-licenseFile.reloadInterval` duration  
 Interval for reloading the license file specified via `-licenseFile`. See <https://victoriametrics.com/products/enterprise/> . This flag is ava

`-loggerDisableTimestamps`  
 Whether to disable writing timestamps in logs

`-loggerErrorsPerSecondLimit` int  
 Per-second limit on the number of ERROR messages. If more than the given number of errors are emitted per second, the remaining errors are :

`-loggerFormat` string  
 Format for logs. Possible values: default, json (default "default")

`-loggerJSONFields` string  
 Allows renaming fields in JSON formatted logs. Example: "ts:timestamp,msg:message" renames "ts" to "timestamp" and "msg" to "message". Supp

`-loggerLevel` string  
 Minimum level of errors to log. Possible values: INFO, WARN, ERROR, FATAL, PANIC (default "INFO")

`-loggerMaxArgLen` int  
 The maximum length of a single logged argument. Longer arguments are replaced with 'arg\_start..arg\_end', where 'arg\_start' and 'arg\_end' is

`-loggerOutput` string  
 Output for the logs. Supported values: stderr, stdout (default "stderr")

`-loggerTimezone` string  
 Timezone to use for timestamps in logs. Timezone must be a valid IANA Time Zone. For example: America/New\_York, Europe/Berlin, Etc/GMT+3 or

`-loggerWarnsPerSecondLimit` int  
 Per-second limit on the number of WARN messages. If more than the given number of warns are emitted per second, then the remaining warns are

`-memory.allowedBytes` size  
 Allowed size of system memory VictoriaMetrics caches may occupy. This option overrides `-memory.allowedPercent` if set to a non-zero value. Tr  
 Supports the following optional suffixes for size values: KB, MB, GB, TB, KiB, MiB, GiB, TiB (default 0)

`-memory.allowedPercent` float  
 Allowed percent of system memory VictoriaMetrics caches may occupy. See also `-memory.allowedBytes`. Too low a value may increase cache miss

`-metrics.exposeMetadata`  
 Whether to expose TYPE and HELP metadata at the `/metrics` page, which is exposed at `-httpListenAddr` . The metadata may be needed when the `/m`

`-metricsAuthKey` value  
 Auth key for `/metrics` endpoint. It must be passed via `authKey` query arg. It overrides `-httpAuth.*`  
 Flag value can be read from the given file when using `-metricsAuthKey=file:///abs/path/to/file` or `-metricsAuthKey=file:///relative/path/to,`

`-mtls` array  
 Whether to require valid client certificate for https requests to the corresponding `-httpListenAddr` . This flag works only if `-tls` flag is :  
 Supports array of values separated by comma or specified via multiple flags.  
 Empty values are set to false.

`-mtlsCAFile` array  
 Optional path to TLS Root CA for verifying client certificates at the corresponding `-httpListenAddr` when `-mtls` is enabled. By default the h  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.basicAuth.password` array  
 Optional basic auth password for `-notifier.url`  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.basicAuth.passwordFile` array  
 Optional path to basic auth password file for `-notifier.url`  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.basicAuth.username` array  
 Optional basic auth username for `-notifier.url`  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.bearerToken` array  
 Optional bearer token for `-notifier.url`  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.bearerTokenFile` array  
 Optional path to bearer token file for `-notifier.url`  
 Supports an array of values separated by comma or specified via multiple flags.  
 Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.blackhole`  
 Whether to blackhole alerting notifications. Enable this flag if you want vmalert to evaluate alerting rules without sending any notificati

`-notifier.config` string  
 Path to configuration file for notifiers

`-notifier.headers` array

Optional HTTP headers to send with each request to the corresponding `-notifier.url`. For example, `-remoteWrite.headers='My-Auth:foobar'` would support an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.clientID` array  
Optional OAuth2 clientID to use for `-notifier.url`. If multiple args are set, then they are applied independently for the corresponding `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.clientSecret` array  
Optional OAuth2 clientSecret to use for `-notifier.url`. If multiple args are set, then they are applied independently for the corresponding `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.clientSecretFile` array  
Optional OAuth2 clientSecretFile to use for `-notifier.url`. If multiple args are set, then they are applied independently for the corresponding `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.endpointParams` array  
Optional OAuth2 endpoint parameters to use for the corresponding `-notifier.url`. The endpoint parameters must be set in JSON format: {"param": "value"}. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.scopes` array  
Optional OAuth2 scopes to use for `-notifier.url`. Scopes must be delimited by ';'. If multiple args are set, then they are applied independently for the corresponding `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.oauth2.tokenUrl` array  
Optional OAuth2 tokenURL to use for `-notifier.url`. If multiple args are set, then they are applied independently for the corresponding `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.sendTimeout` array  
Timeout when sending alerts to the corresponding `-notifier.url` (default 10s). Supports array of values separated by comma or specified via multiple flags. Empty values are set to default value.

`-notifier.showURL`  
Whether to avoid stripping sensitive information such as passwords from URL in log messages or UI for `-notifier.url`. It is hidden by default.

`-notifier.suppressDuplicateTargetErrors`  
Whether to suppress 'duplicate target' errors during discovery.

`-notifier.tlsCAFile` array  
Optional path to TLS CA file to use for verifying connections to `-notifier.url`. By default, system CA is used. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.tlsCertFile` array  
Optional path to client-side TLS certificate file to use when connecting to `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.tlsInsecureSkipVerify` array  
Whether to skip TLS verification when connecting to `-notifier.url`. Supports array of values separated by comma or specified via multiple flags. Empty values are set to false.

`-notifier.tlsKeyFile` array  
Optional path to client-side TLS certificate key to use when connecting to `-notifier.url`. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.tlsServerName` array  
Optional TLS server name to use for connections to `-notifier.url`. By default, the server name from `-notifier.url` is used. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-notifier.url` array  
Prometheus Alertmanager URL, e.g. `http://127.0.0.1:9093`. List all Alertmanager URLs if it runs in the cluster mode to ensure high availability. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.

`-pprofAuthKey` value  
Auth key for `/debug/pprof/*` endpoints. It must be passed via `authKey` query arg. It overrides `-httpAuth.*`. Flag value can be read from the given file when using `-pprofAuthKey=file:///abs/path/to/file` or `-pprofAuthKey=file:///relative/path/to/file`.

`-promscrape.consul.waitTime` duration  
Wait time used by Consul service discovery. Default value is used if not set.

`-promscrape.consulSDCheckInterval` duration  
Interval for checking for changes in Consul. This works only if `consul_sd_configs` is configured in `'-promscrape.config'` file. See <https://docs.victoriametrics.com/victoriametrics/promscrape/#consul>.

`-promscrape.discovery.concurrency` int  
The maximum number of concurrent requests to Prometheus autodiscovery API (Consul, Kubernetes, etc.) (default 100).

`-promscrape.discovery.concurrentWaitTime` duration  
The maximum duration for waiting to perform API requests if more than `-promscrape.discovery.concurrency` requests are simultaneously performed.

`-promscrape.dnsSDCheckInterval` duration  
Interval for checking for changes in DNS. This works only if `dns_sd_configs` is configured in `'-promscrape.config'` file. See <https://docs.victoriametrics.com/victoriametrics/promscrape/#dns>.

`-pushmetrics.disableCompression`  
Whether to disable request body compression when pushing metrics to every `-pushmetrics.url`.

`-pushmetrics.extraLabel` array  
Optional labels to add to metrics pushed to every `-pushmetrics.url`. For example, `-pushmetrics.extraLabel='instance="foo"'` adds `instance="foo"` label. Supports an array of values separated by comma or specified via multiple flags. Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.



```

-pushmetrics.header array
  Optional HTTP request header to send to every -pushmetrics.url . For example, -pushmetrics.header='Authorization: Basic foobar' adds 'Autho
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-pushmetrics.interval duration
  Interval for pushing metrics to every -pushmetrics.url (default 10s)
-pushmetrics.url array
  Optional URL to push metrics exposed at /metrics page. See https://docs.victoriametrics.com/victoriametrics/single-server-victoriametrics/#
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-reloadAuthKey value
  Auth key for /-/reload http endpoint. It must be passed via authKey query arg. It overrides -httpAuth.*
  Flag value can be read from the given file when using -reloadAuthKey=file:///abs/path/to/file or -reloadAuthKey=file:///./relative/path/to/fi
-remoteRead.basicAuth.password string
  Optional basic auth password for -remoteRead.url
-remoteRead.basicAuth.passwordFile string
  Optional path to basic auth password to use for -remoteRead.url
-remoteRead.basicAuth.username string
  Optional basic auth username for -remoteRead.url
-remoteRead.bearerToken string
  Optional bearer auth token to use for -remoteRead.url.
-remoteRead.bearerTokenFile string
  Optional path to bearer token file to use for -remoteRead.url.
-remoteRead.disablePathAppend
  Whether to disable automatic appending of '/api/v1/query' or '/select/logsql/stats_query' path to the configured -datasource.url and -remot
-remoteRead.headers string
  Optional HTTP headers to send with each request to the corresponding -remoteRead.url. For example, -remoteRead.headers='My-Auth:foobar' wou
-remoteRead.idleConnTimeout duration
  Defines a duration for idle (keep-alive connections) to exist. Consider settings this value less to the value of "-http.idleConnTimeout". I
-remoteRead.lookback duration
  Lookback defines how far to look into past for alerts timeseries. For example, if lookback=1h then range from now() to now()-1h will be sca
-remoteRead.oauth2.clientID string
  Optional OAuth2 clientID to use for -remoteRead.url.
-remoteRead.oauth2.clientSecret string
  Optional OAuth2 clientSecret to use for -remoteRead.url.
-remoteRead.oauth2.clientSecretFile string
  Optional OAuth2 clientSecretFile to use for -remoteRead.url.
-remoteRead.oauth2.endpointParams string
  Optional OAuth2 endpoint parameters to use for -remoteRead.url . The endpoint parameters must be set in JSON format: {"param1":"value1",...
-remoteRead.oauth2.scopes string
  Optional OAuth2 scopes to use for -remoteRead.url. Scopes must be delimited by ';'.
-remoteRead.oauth2.tokenUrl string
  Optional OAuth2 tokenURL to use for -remoteRead.url.
-remoteRead.showURL
  Whether to show -remoteRead.url in the exported metrics. It is hidden by default, since it can contain sensitive info such as auth key
-remoteRead.tlsCAFile string
  Optional path to TLS CA file to use for verifying connections to -remoteRead.url. By default, system CA is used
-remoteRead.tlsCertFile string
  Optional path to client-side TLS certificate file to use when connecting to -remoteRead.url
-remoteRead.tlsInsecureSkipVerify
  Whether to skip tls verification when connecting to -remoteRead.url
-remoteRead.tlsKeyFile string
  Optional path to client-side TLS certificate key to use when connecting to -remoteRead.url
-remoteRead.tlsServerName string
  Optional TLS server name to use for connections to -remoteRead.url. By default, the server name from -remoteRead.url is used
-remoteRead.url vmalert
  Optional URL to datasource compatible with MetricsQL. It can be single node VictoriaMetrics or vmselect.Remote read is used to restore aler
-remoteWrite.basicAuth.password string
  Optional basic auth password for -remoteWrite.url
-remoteWrite.basicAuth.passwordFile string
  Optional path to basic auth password to use for -remoteWrite.url
-remoteWrite.basicAuth.username string
  Optional basic auth username for -remoteWrite.url
-remoteWrite.bearerToken string
  Optional bearer auth token to use for -remoteWrite.url.
-remoteWrite.bearerTokenFile string
  Optional path to bearer token file to use for -remoteWrite.url.
-remoteWrite.concurrency int
  Defines number of writers for concurrent writing into remote write endpoint. Default value depends on the number of available CPU cores. (d
-remoteWrite.disablePathAppend
  Whether to disable automatic appending of '/api/v1/write' path to the configured -remoteWrite.url.
-remoteWrite.flushInterval duration
  Defines interval of flushes to remote write endpoint (default 2s)
-remoteWrite.headers string
  Optional HTTP headers to send with each request to the corresponding -remoteWrite.url. For example, -remoteWrite.headers='My-Auth:foobar' w
-remoteWrite.idleConnTimeout duration
  Defines a duration for idle (keep-alive connections) to exist. Consider settings this value less to the value of "-http.idleConnTimeout". I
-remoteWrite.maxBatchSize int

```

Defines max number of timeseries to be flushed at once (default 10000)

`-remoteWrite.maxQueueSize` int  
Defines the max number of pending datapoints to remote write endpoint (default 100000)

`-remoteWrite.oauth2.clientID` string  
Optional OAuth2 clientID to use for `-remoteWrite.url`

`-remoteWrite.oauth2.clientSecret` string  
Optional OAuth2 clientSecret to use for `-remoteWrite.url`

`-remoteWrite.oauth2.clientSecretFile` string  
Optional OAuth2 clientSecretFile to use for `-remoteWrite.url`

`-remoteWrite.oauth2.endpointParams` string  
Optional OAuth2 endpoint parameters to use for `-remoteWrite.url`. The endpoint parameters must be set in JSON format: `{"param1":"value1",...`

`-remoteWrite.oauth2.scopes` string  
Optional OAuth2 scopes to use for `-notifier.url`. Scopes must be delimited by `';`.

`-remoteWrite.oauth2.tokenUrl` string  
Optional OAuth2 tokenURL to use for `-notifier.url`.

`-remoteWrite.retryMaxTime` duration  
The max time spent on retry attempts for the failed remote-write request. Change this value if it is expected for `remoteWrite.url` to be unreachable.

`-remoteWrite.retryMinInterval` duration  
The minimum delay between retry attempts. Every next retry attempt will double the delay to prevent hammering of remote database. See also `-replay.rulesDelay`

`-remoteWrite.sendTimeout` duration  
Timeout for sending data to the configured `-remoteWrite.url`. (default 30s)

`-remoteWrite.showURL`  
Whether to show `-remoteWrite.url` in the exported metrics. It is hidden by default, since it can contain sensitive info such as auth key

`-remoteWrite.tlsCAFile` string  
Optional path to TLS CA file to use for verifying connections to `-remoteWrite.url`. By default, system CA is used

`-remoteWrite.tlsCertFile` string  
Optional path to client-side TLS certificate file to use when connecting to `-remoteWrite.url`

`-remoteWrite.tlsInsecureSkipVerify`  
Whether to skip tls verification when connecting to `-remoteWrite.url`

`-remoteWrite.tlsKeyFile` string  
Optional path to client-side TLS certificate key to use when connecting to `-remoteWrite.url`

`-remoteWrite.tlsServerName` string  
Optional TLS server name to use for connections to `-remoteWrite.url`. By default, the server name from `-remoteWrite.url` is used

`-remoteWrite.url` string  
Optional URL to VictoriaMetrics or vminsert where to persist alerts state and recording rules results in form of timeseries. Supports address and authentication parameters.

`-replay.disableProgressBar`  
Whether to disable rendering progress bars during the replay. Progress bar rendering might be verbose or break the logs parsing, so it is recommended to disable it.

`-replay.maxDatapointsPerQuery` int  
Max number of data points expected in one request. It affects the max time range for every `'/query_range'` request during the replay. The higher the value, the more data points are requested in one query, which may lead to increased response time.

`-replay.ruleRetryAttempts` int  
Defines how many retries to make before giving up on rule if request for it returns an error. (default 5)

`-replay.rulesDelay` duration  
Delay between rules evaluation within the group. Could be important if there are chained rules inside the group and processing need to wait for previous rule results.

`-replay.timeFrom` string  
The time filter in RFC3339 format to start the replay from. E.g. `'2020-01-01T20:07:00Z'`

`-replay.timeTo` string  
The time filter in RFC3339 format to finish the replay by. E.g. `'2020-01-01T20:07:00Z'`. By default, is set to the current time.

`-rule` array  
Path to the files or http url with alerting and/or recording rules in YAML format.  
Supports hierarchical patterns and regexes.  
Examples:  
`-rule="/path/to/file"`. Path to a single file with alerting rules.  
`-rule="http://<some-server-addr>/path/to/rules"`. HTTP URL to a page with alerting rules.  
`-rule="dir/*.yaml"` `-rule="/*.yaml"` `-rule="gcs://vmalert-rules/tenant_{TENANT_ID}/prod"`.  
`-rule="dir/**/*.yaml"`. Includes all the .yaml files in "dir" subfolders recursively.  
Rule files support YAML multi-document. Files may contain `%(ENV_VAR)` placeholders, which are substituted by the corresponding env vars.

Enterprise version of vmalert supports S3 and GCS paths to rules.  
For example: `gs://bucket/path/to/rules`, `s3://bucket/path/to/rules`  
S3 and GCS paths support only matching by prefix, e.g. `s3://bucket/dir/rule_ matches`  
all files with prefix `rule_` in folder `dir`.  
See <https://docs.victoriametrics.com/victoriametrics/vmalert/#reading-rules-from-object-storage>

Supports an array of values separated by comma or specified via multiple flags.  
Value can contain comma inside single-quoted or double-quoted string, `{}`, `[]` and `()` braces.

`-rule.defaultRuleType` string  
Default type for rule expressions, can be overridden via `"type"` parameter on the group level, see <https://docs.victoriametrics.com/victoriametrics/vmalert/#rule-type>

`-rule.evalDelay` duration  
Adjustment of the `'time'` parameter for rule evaluation requests to compensate intentional data delay from the datasource. Normally, should be 0.

`-rule.maxResolveDuration` duration  
Limits the maximum duration for automatic alert expiration, which by default is 4 times `evaluationInterval` of the parent group

`-rule.resendDelay` duration  
Minimum amount of time to wait before resending an alert to notifier.

`-rule.stripFilePath`  
Whether to strip file path in responses from the `api/v1/rules` API for files configured via `-rule` cmd-line flag. For example, the file path `path/to/file.yaml` will be replaced by `rule_`.

`-rule.templates` array  
Path or glob pattern to location with go template definitions for rules annotations templating. Flag can be specified multiple times.  
Examples:

```

-rule.templates="/path/to/file". Path to a single file with go templates
-rule.templates="dir/*.tpl" -rule.templates="/*.tpl". Relative path to all .tpl files in "dir" folder,
absolute path to all .tpl files in root.
-rule.templates="dir/**/*.tpl". Includes all the .tpl files in "dir" subfolders recursively.

Supports an array of values separated by comma or specified via multiple flags.
Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-rule.updateEntriesLimit int
  Defines the max number of rule's state updates stored in-memory. Rule's updates are available on rule's Details page and are used for debug
-rule.validateExpressions
  Whether to validate rules expressions for different types. (default true)
-rule.validateTemplates
  Whether to validate annotation and label templates (default true)
-s3.configFilePath string
  Path to file with S3 configs. Configs are loaded from default location if not set.
  See https://docs.aws.amazon.com/general/latest/gr/aws-security-credentials.html . This flag is available only in Enterprise binaries. See h
-s3.configProfile string
  Profile name for S3 configs. If no set, the value of the environment variable will be loaded (AWS_PROFILE or AWS_DEFAULT_PROFILE), or if bo
-s3.credsFilePath string
  Path to file with GCS or S3 credentials. Credentials are loaded from default locations if not set.
  See https://cloud.google.com/iam/docs/creating-managing-service-account-keys and https://docs.aws.amazon.com/general/latest/gr/aws-security
-s3.customEndpoint string
  Custom S3 endpoint for use with S3-compatible storages (e.g. MinIO). S3 is used if not set. This flag is available only in Enterprise binar
-s3.forcePathStyle
  Prefixing endpoint with bucket name when set false, true by default. This flag is available only in Enterprise binaries. See https://docs.v
-tls array
  Whether to enable TLS for incoming HTTP requests at the given -httpListenAddr (aka https). -tlsCertFile and -tlsKeyFile must be set if -tls
  Supports array of values separated by comma or specified via multiple flags.
  Empty values are set to false.
-tlsAutocertCacheDir string
  Directory to store TLS certificates issued via Let's Encrypt. Certificates are lost on restarts if this flag isn't set. This flag is availa
-tlsAutocertEmail string
  Contact email for the issued Let's Encrypt TLS certificates. See also -tlsAutocertHosts and -tlsAutocertCacheDir .This flag is available on
-tlsAutocertHosts array
  Optional hostnames for automatic issuing of Let's Encrypt TLS certificates. These hostnames must be reachable at -httpListenAddr . The -htt
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-tlsCertFile array
  Path to file with TLS certificate for the corresponding -httpListenAddr if -tls is set. Prefer ECDSA certs instead of RSA certs as RSA cert:
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-tlsCipherSuites array
  Optional list of TLS cipher suites for incoming requests over HTTPS if -tls is set. See the list of supported cipher suites at https://pkg.
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-tlsKeyFile array
  Path to file with TLS key for the corresponding -httpListenAddr if -tls is set. The provided key file is automatically re-read every second
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-tlsMinVersion array
  Optional minimum TLS version to use for the corresponding -httpListenAddr if -tls is set. Supported values: TLS10, TLS11, TLS12, TLS13
  Supports an array of values separated by comma or specified via multiple flags.
  Value can contain comma inside single-quoted or double-quoted string, {}, [] and () braces.
-version
  Show VictoriaMetrics version

```

## Hot config reload

vmaalert supports "hot" config reload via the following methods:

- send SIGHUP signal to vmaalert process;
- send GET request to `/-/reload` endpoint (this endpoint can be protected with `-reloadAuthKey` command-line flag);
- configure `-configCheckInterval` flag for periodic reload on config change.

## URL params

To set additional URL params for `datasource.url`, `remoteWrite.url` or `remoteRead.url` just add them in address: `-datasource.url=http://localhost:8428?nocache=1`.

To set additional URL params for specific [group of rules](#) modify the `params` group:



```
groups:
- name: TestGroup
  params:
    denyPartialResponse: ["true"]
    extra_label: ["env=dev"]
```

Please note, `params` are used only for executing rules expressions (requests to `datasource.url`). If there would be a conflict between URL params set in `datasource.url` flag and params in group definition the latter will have higher priority.

## Chaining groups

For chaining groups, they must be executed in a specific order, and the next group should be executed after the results from previous group are available in the datasource. In `vmalert`, user can specify `eval_offset` to achieve that [?](#).

For example:



```
groups:
- name: BaseGroup
  interval: 1m
  eval_offset: 10s
  rules:
    - record: http_server_request_duration_seconds:sum_rate:5m:http_get
      expr: |
        sum without(instance, pod) (
          rate(
            http_server_request_duration_seconds{
              http_request_method="GET"
            }[5m]
          )
        )
    - record: http_server_request_duration_seconds:sum_rate:5m:http_post
      expr: |
        sum without(instance, pod) (
          rate(
            http_server_request_duration_seconds{
              http_request_method="POST"
            }[5m]
          )
        )
- name: TopGroup
  interval: 1m
  eval_offset: 40s
  rules:
    - record: http_server_request_duration_seconds:sum_rate:5m:merged
      expr: |
        http_server_request_duration_seconds:sum_rate:5m:http_get
        or
        http_server_request_duration_seconds:sum_rate:5m:http_post
```

This configuration ensures that rules in `BaseGroup` are executed at (assuming `vmalert` starts at `12:00:00`):



```
[12:00:10, 12:01:10, 12:02:10, 12:03:10...]
```

while rules in group `TopGroup` are executed at:



```
[12:00:40, 12:01:40, 12:02:40, 12:03:40...]
```

As a result, TopGroup always gets the latest results of BaseGroup .

By default, the `eval_offset` values should be at least 30 seconds apart to accommodate the `-search.latencyOffset(default 30s)` command-line flag at vmselect or VictoriaMetrics single-node. The minimum `eval_offset` gap can be adjusted accordingly with `-search.latencyOffset` .

## Notifier configuration file

Notifier also supports configuration via file specified with flag `notifier.config` :



```
./bin/vmalert -rule=app/vmalert/config/testdata/rules.good.rules \
-datasource.url=http://localhost:8428 \
-notifier.config=app/vmalert/notifier/testdata/consul.good.yaml
```

The configuration file allows to configure static notifiers, discover notifiers via [Consul](#) and [DNS](#) : For example:



```
static_configs:
- targets:
  # support using full url
  - 'http://alertmanager:9093/test/api/v2/alerts'
  - 'https://alertmanager:9093/api/v2/alerts'
  # the following target with only host:port will be used as <scheme>://localhost:9093/<path_prefix>/api/v2/alerts
  - localhost:9093

consul_sd_configs:
- server: localhost:8500
  services:
  - alertmanager

dns_sd_configs:
- names:
  - my.domain.com
  type: 'A'
  port: 9093
```

The list of configured or discovered Notifiers can be explored via [UI](#) . If Alertmanager runs in cluster mode then all its URLs needs to be available during discovery to ensure [high availability](#) .

The configuration file [specification](#) is the following:



```
# Per-target Notifier timeout when pushing alerts.
[ timeout: <duration> | default = 10s ]

# Prefix for the HTTP path alerts are pushed to.
[ path_prefix: <path> | default = / ]

# Configures the protocol scheme used for requests.
[ scheme: <scheme> | default = http ]

# Sets the `Authorization` header on every request with the
# configured username and password.
# password and password_file are mutually exclusive.
basic_auth:
[ username: <string> ]
[ password: <string> ]
[ password_file: <string> ]

# Optional `Authorization` header configuration.
authorization:
# Sets the authentication type.
```

```

[ type: <string> | default: Bearer ]
# Sets the credentials. It is mutually exclusive with
# `credentials_file`.
[ credentials: <secret> ]
# Sets the credentials to the credentials read from the configured file.
# It is mutually exclusive with `credentials`.
[ credentials_file: <filename> ]

# Configures the scrape request's TLS settings.
# see https://prometheus.io/docs/prometheus/latest/configuration/configuration/#tls_config
tls_config:
[ <tls_config> ]

# Configures Bearer authentication token via string
bearer_token: <string>
# or by passing path to the file with token.
bearer_token_file: <string>

# Configures OAuth 2.0 authentication
# see https://prometheus.io/docs/prometheus/latest/configuration/configuration/#oauth2
oauth2:
[ <oauth2_config> ]

# Optional list of HTTP headers in form `header-name: value`
# applied for all requests to notifiers
# For example:
# headers:
#   - "CustomHeader: foo"
#   - "CustomHeader2: bar"
headers:
[ <string>, ... ]

# List of labeled statically configured Notifiers.
#
# Each list of targets may be additionally instructed with
# authorization params. Target's authorization params will
# inherit params from global authorization params if there
# are no conflicts.
static_configs:
[ - targets: ]
  [ - '<host>' ]
  [ oauth2 ]
  [ basic_auth ]
  [ authorization ]
  [ tls_config ]
  [ bearer_token ]
  [ bearer_token_file ]
  [ headers ]

# List of Consul service discovery configurations.
# See https://prometheus.io/docs/prometheus/latest/configuration/configuration/#consul_sd_config
consul_sd_configs:
[ - <consul_sd_config> ... ]

# List of DNS service discovery configurations.
# See https://prometheus.io/docs/prometheus/latest/configuration/configuration/#dns_sd_config
dns_sd_configs:
[ - <dns_sd_config> ... ]

# List of relabel configurations for entities discovered via service discovery.
# Supports the same relabeling features as the rest of VictoriaMetrics components.
# See https://docs.victoriametrics.com/victoriametrics/vmagent/#relabeling
relabel_configs:
[ - <relabel_config> ... ]

# List of relabel configurations for alert labels sent via Notifier.
# Supports the same relabeling features as the rest of VictoriaMetrics components.
# See https://docs.victoriametrics.com/victoriametrics/vmagent/#relabeling
alert_relabel_configs:
[ - <relabel_config> ... ]

```

The configuration file can be [hot-reloaded](#).

## Contributing

vmaAlert is mostly designed and built by VictoriaMetrics community. Feel free to share your experience and ideas for improving this software. Please keep simplicity as the main priority.

## How to build from sources

It is recommended using [binary releases](#)

- vmaAlert is located in vmutils-\* archives there.

## Docker image

You can build vmaAlert docker image from source and push it to your own docker repository. Run the following commands from the root folder of [the repository](#):



```
make package-vmaAlert
docker tag victoria-metrics/vmaAlert:version my-repo:my-version-name
docker push my-repo:my-version-name
```

To run the built image in victoria-metrics-k8s-stack or VMAAlert CR object apply the following config change:



```
kind: VMAAlert
spec:
  image:
    repository: my-repo
    tag: my-version-name
```

## Development build

1. [Install Go](#).
2. Run `make vmaAlert` from the root folder of [the repository](#). It builds vmaAlert binary and puts it into the `bin` folder.

## Production build

1. [Install docker](#).
2. Run `make vmaAlert-prod` from the root folder of [the repository](#). It builds vmaAlert-prod binary and puts it into the `bin` folder.

## ARM build

ARM build may run on Raspberry Pi or on [energy-efficient ARM servers](#).

## Development ARM build

1. [Install Go](#).
2. Run `make vmaAlert-linux-arm` or `make vmaAlert-linux-arm64` from the root folder of [the repository](#). It builds vmaAlert-linux-arm or vmaAlert-linux-arm64 binary respectively and puts it into the `bin` folder.

## Production ARM build

1. [Install docker](#).

2. Run `make vmaalert-linux-arm-prod` or `make vmaalert-linux-arm64-prod` from the root folder of [the repository](#). It builds `vmaalert-linux-arm-prod` or `vmaalert-linux-arm64-prod` binary respectively and puts it into the `bin` folder.

[Previous](#)

vmagent

[Next](#)

vmauth